

Predictive Modeling Homework 3

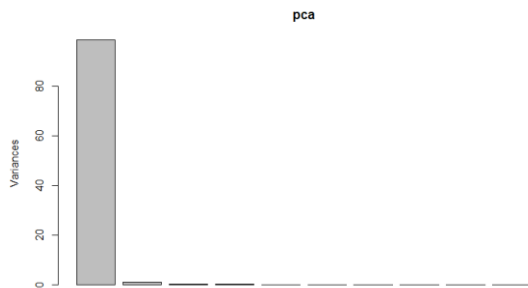
Ian Boulis

6.1

A. Start R and use these commands to load the data:

```
> library(caret)
> data(tecator)
```

B. Use PCA to determine the effective dimension of these data, use 95% as a cut-off value.



From the screeplot, it appears that almost all the variance is accounted for in the first PCA. This is verified by computing the variance of the PCAs:

```
> head(varpca)
[1] 98.626192582  0.969705229  0.279324276  0.114429868  0.006460911  0.002624591
```

C. Split the data into a training and a test set (80/20). Use 3-fold CV to build the PCR model (95% cut off). What is the average RMSE?

Linear Regression

174 samples
100 predictors

No pre-processing
Resampling: Cross-validated (3 fold)
Summary of sample sizes: 117, 115, 116
Resampling results:

RMSE	Rsquared	MAE
7.2135	0.7921591	4.190897

D. Use the model on (C) to predict the response for the test set. What is the RMSE?

Using the model from part C to predict on the test set resulted in a RMSE of 2.8921448 and R^2 value of 0.94

```
> defaultsummary(dapa)
      RMSE  Rsquared    MAE
2.8921448 0.9472097 1.7976082
```

6.2

B. Filter out the predictors that have low frequency using the nearZeroVar function. How many predictors are left for modeling?

values	
nzv	int [1:719] 7 8 9 10 13 14 17 18 19 22 ...

There is a total of 719 predictors with a zero or near zero variance. From our 1,107 predictors, we have 388 predictors left for model building.

C. Split the data into a training and test set, pre-process the data and tune a PLS model. How many latent variables are optimal and what is the corresponding resampled estimate of R^2 ?

Partial Least Squares

133 samples
388 predictors

Pre-processing: centered (388), scaled (388)

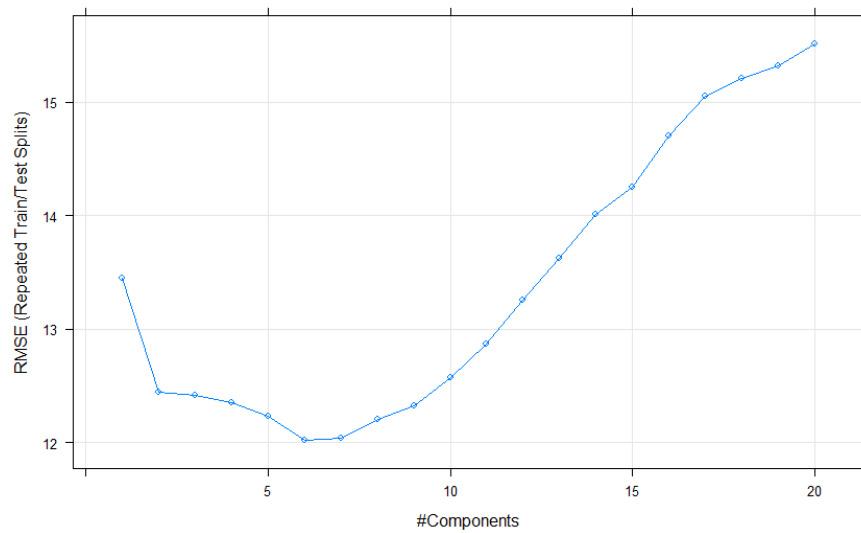
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 101, 101, 101, 101, 101, 101, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	MAE
1	13.44745	0.3158445	10.296566
2	12.44250	0.4304343	8.625833
3	12.41566	0.4342389	9.140396
4	12.35435	0.4472174	9.125192
5	12.22983	0.4706266	8.956198
6	12.01679	0.4883515	8.829465
7	12.03354	0.4921005	9.012227
8	12.20621	0.4864443	9.259302
9	12.32250	0.4868289	9.247622
10	12.56818	0.4756307	9.411210
11	12.87080	0.4593575	9.605767
12	13.25810	0.4389713	9.838692
13	13.62212	0.4222968	10.099519
14	14.01455	0.4023258	10.352074
15	14.25061	0.3935322	10.521174
16	14.70161	0.3712271	10.869507
17	15.04990	0.3572226	11.100422
18	15.21098	0.3536296	11.234272
19	15.31762	0.3520845	11.388072
20	15.51716	0.3471868	11.555851

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was ncomp = 6.



As seen in the pls output from R, the optimal number of latent variables is 6, with a corresponding value R^2 of 0.4883515. However, the number of components with the highest R^2 value is 7 components, with an R^2 value of 0.4921005.

D. Predict the response for the test set. What is the test set estimate of R^2 ?

Using the same model from part (C) resulted in a RMSE of 10.3124854 and an R^2 value of 0.70

```
> defaultSummary(dapa)
      RMSE  Rsquared    MAE
10.3124854 0.7055006  7.7924576
```

6.3

B. A small percentage of calls in the predictor set contains missing values. Use an imputation function to fill in these missing values

```
impute <- preProcess(trp, method=c("center", "scale", "BoxCox", "knnImpute"))
```

After splitting the data, I used the preProcess function on the predictor training set to compact the pre-processing steps. I took a look at the distribution for the response variable (Yield in this data set), decided to center and scale then preform the Box-Cox transformation, then used the built in knnImpute method to impute the data. The default value for k in this context is k=5.

C. Split the data into a training and a test set (80%/20% in training/test sets), pre-process the data, and use 3-fold CV to build lm, Ridge, lasso, and elastic net models on the training set. What is the average RMSE for each model on the training set? What are the best tuning parameters for each model?

As mentioned in part B, the data was prepressed with center and scaling, and then a Box-Cox transformation.

Model	Best Tuning Parameter	Training		Testing	
		RMSE	R ²	RMSE	R ²
Linear Model		2.714181	0.4164144	14.2545	0.04543042
Ridge	$\Lambda = 0.1$	3.60031	0.3977611	1.980587	0.3005124
LASSO	Fraction= 0.1	1.122202	0.657	1.671107	0.15535395
E-Net	Fraction=0.525 $\Lambda = 0.1$	1.135579	0.6296772	1.297186	0.5380803

Linear Regression

144 samples
46 predictor

No pre-processing
Resampling: Cross-validated (3 fold)
Summary of sample sizes: 96, 96, 96
Resampling results:

RMSE	Rsquared	MAE
2.714181	0.4164144	1.371673

Tuning parameter 'intercept' was held constant at a value of TRUE

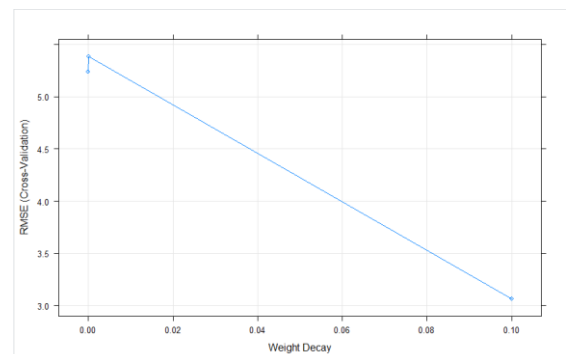
Ridge Regression

144 samples
46 predictor

No pre-processing
Resampling: Cross-validated (3 fold)
Summary of sample sizes: 96, 96, 96
Resampling results across tuning parameters:

lambda	RMSE	Rsquared	MAE
0e+00	9.506434	0.1129453	2.367087
1e-04	9.469312	0.1178802	2.355339
1e-01	3.600305	0.3977611	1.351270

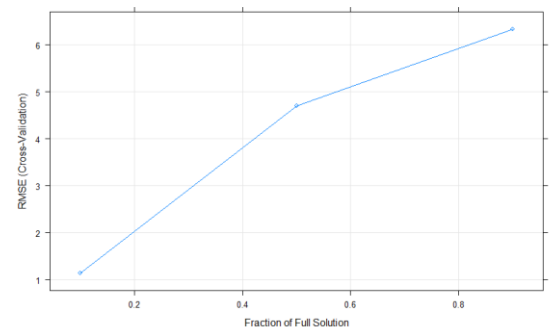
RMSE was used to select the optimal model using the smallest value.
The final value used for the model was lambda = 0.1.



The lasso
 144 samples
 46 predictor
 No pre-processing
 Resampling: Cross-Validated (3 fold)
 Summary of sample sizes: 96, 96, 96
 Resampling results across tuning parameters:

fraction	RMSE	Rsquared	MAE
0.1	1.122202	0.6570048	0.9113331
0.5	3.393380	0.2721675	1.4136536
0.9	4.616505	0.1955870	1.7368916

RMSE was used to select the optimal model using the smallest value.
 The final value used for the model was fraction = 0.1.

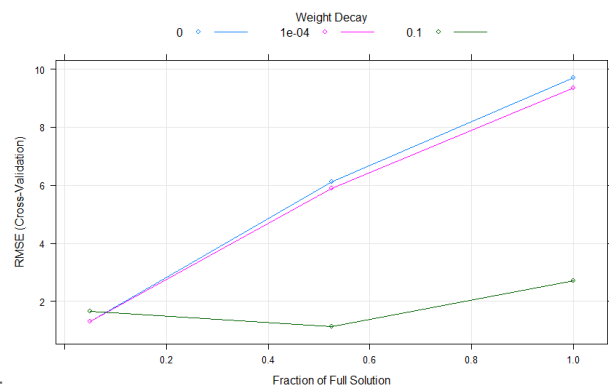


Elasticnet

144 samples
 46 predictor
 No pre-processing
 Resampling: Cross-Validated (3 fold)
 Summary of sample sizes: 96, 96, 96
 Resampling results across tuning parameters:

lambda	fraction	RMSE	Rsquared	MAE
0e+00	0.050	1.298305	0.6186981	1.0442678
0e+00	0.525	6.121063	0.2764137	1.8889712
0e+00	1.000	9.705170	0.1606099	2.5893353
1e-04	0.050	1.315788	0.6174787	1.0572792
1e-04	0.525	5.887176	0.2843516	1.8437054
1e-04	1.000	9.357351	0.1638396	2.5133665
1e-01	0.050	1.663812	0.4833315	1.3337512
1e-01	0.525	1.135579	0.6296772	0.9127071
1e-01	1.000	2.719747	0.4157516	1.3019369

RMSE was used to select the optimal model using the smallest value.
 The final values used for the model were fraction = 0.525 and lambda = 0.1.



D. Predict the response for the test set for each model in (C). What is the RMSE for each model? How does the RMSE compare with the RMSE in the training set?

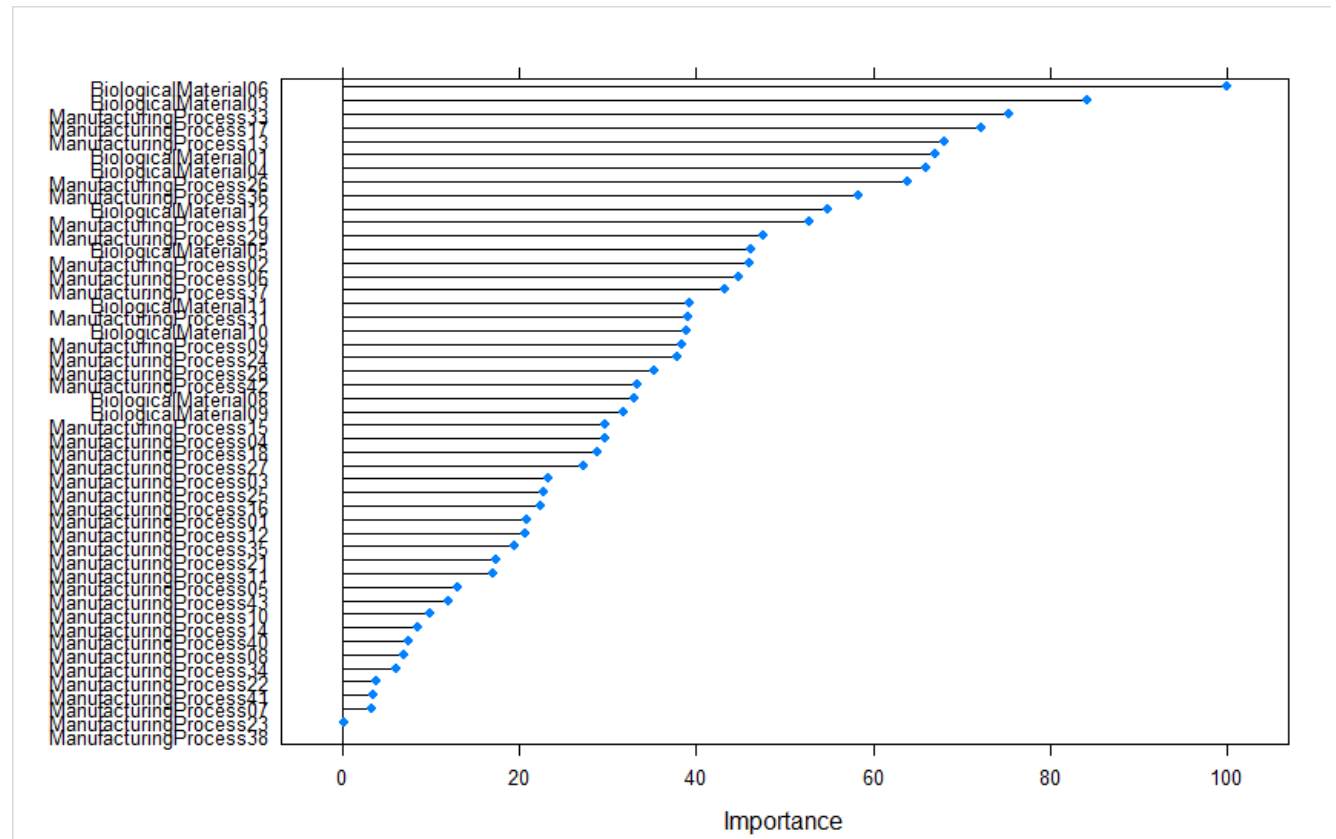
The root mean squared error along with the R^2 values for all four of the models can be seen in the table in part C. For the linear model, there was quite a stark difference between the training and test set. The RMSE increased significantly and the R^2 value got significantly worse. For the ridge model, the RMSE showed a better value, with it almost halving between the training and test sets, while the R^2 value did decrease it was not as significant as the RMSE. The lasso model did not exhibit that much of a change in RMSE between training and testing, however, the R^2 did get significantly worse. Finally, the E-Net model exhibited minor changes to both RMSE and R^2 .

E. Which model has the best predictive ability? Explain which model you would use for predicting the response.

I would choose the E-Net model for predicting the response in this data set. From the table in part C, the E-Net- model has the most comparable RMSE and R^2 values. The difference between the training and testing sets for this model are both ~ 0.1 , with RMSE increasing slightly over 0.1, and R^2 decreasing less than 0.1.

F. Based on the best model, which predictors are most important.

Plotting the varImp function we can see that “BiologicalMaterial06” is the most important predictor.



```

##6.1##
##A##
library(AppliedPredictiveModeling)
library(caret)
data(tecator)
##B##
pca <- prcomp(absorp, center=TRUE, scale=TRUE, tol=0.95)
screeplot(pca)
varpca <- pca$sdev^2/sum(pca$sdev^2)*100
head(varpca)
##C##
absorp <- data.frame(absorp)
endpoints <- data.frame(endpoints)
partition <- createDataPartition(endpoints$X2, p=4/5, list=FALSE)
trabsorp <- absorp[partition, ]
teabsorp <- absorp[-partition, ]
trpro <- endpoints$X2[partition]
tepro <- endpoints$X2[-partition ]
ctrl <- trainControl(method="cv", number=3)
Linmod <- train(trabsorp,trpro, method = "lm", trControl = ctrl, preProcess=c("center","scale"))
Linmod
##D##
##test <- train(teabsorp, tepro, method="lm", trControl=ctrl)##
pred <- predict(Linmod, teabsorp)
dapa <- data.frame(obs=tepro, pred=pred)
defaultSummary(dapa)

```



```
##6.2##  
##A##  
library(AppliedPredictiveModeling)  
data("permeability")  
library(pls)  
##B##  
nzs <- nearZeroVar(fingerprints)  
nzv <- fingerprints[, -nzs]  
##C##  
partition <- createDataPartition(permeability, p=4/5, list=FALSE)  
trF <- nzv[partition,]  
trP <- permeability[partition,]  
teF <- nzv[-partition,]  
teP <- permeability[-partition]  
ctrl <- trainControl(method="repeatedcv",3,3)  
pls <- train(trF,trP, method="pls", tuneGrid=expand.grid(ncomp=1:20), trControl=ctrl, preProc =  
c("center", "scale"))  
plot(pls)  
pls  
##D##  
test <- predict(pls,teF)  
dapa <- data.frame(obs=teP, pred=test)  
defaultSummary(dapa)
```

```
##6.3##
```

```
##A##
```

```
library(AppliedPredictiveModeling)
```

```
library(caret)
```

```
library(RANN)
```

```
data("ChemicalManufacturingProcess")
```

```
##B##
```

```
response <- subset(ChemicalManufacturingProcess, select="Yield")
```

```
predict <- subset(ChemicalManufacturingProcess, select=-Yield)
```

```
partition <- createDataPartition(response$Yield, p=4/5, list=FALSE)
```

```
trp <- predict[partition,]
```

```
trr <- response[partition,]
```

```
tep <- predict[-partition,]
```

```
ter <- response[-partition,]
```

```
impute <- preProcess(trp, method=c("center","scale","BoxCox","knnImpute"))
```

```
impute
```

```
##c##D##E##
```

```
trainP <- predict(impute,trp)
```

```
nzv <- nearZeroVar(trainP)
```

```
trainP <- trainP[-nzv]
```

```
corr <- cor(trainP)
```

```
hc <- findCorrelation(corr)
```

```
trainP <- trainP[, -hc]
```

```
timpute <- preProcess(tep, method=c("BoxCox","center","scale","knnImpute"))
```

```
tep <- predict(timpute,tep)
```

```
nzv <- nearZeroVar(tep)
```

```
tep <- tep[-nzv]
```

```
corr <- cor(tep)
```

```
hc <- findCorrelation(corr)
```

```
tep <- tep[, -hc]
```

```
ctrl <- trainControl(method="cv", 3)
```

```
lm <- train(x=trainP, y=trr, method="lm", trControl=ctrl)
```

```
ridge <- train(x=trainP, y=trr, method="ridge", trControl=ctrl)
```

```
lasso <- train(x=trainP, y=trr, method="lasso", trControl=ctrl)
```

```
enet <- train(x=trainP, y=trr, method="enet", trControl=ctrl)
```

```
tlm <- train(x=tep, y=ter, method="lm", trControl=ctrl) ##tlm = test linear model##
```

```
tridge <- train(x=tep, y=ter, method="ridge", trControl=ctrl)
```

```
tlasso <- train(x=tep, y=ter, method="lasso", trControl=ctrl)
```

```
tenet <- train(x=tep, y=ter, method="enet", trControl=ctrl)
```

```
##F##
```

```
plot(varImp(tenet))
```